

NLP

Text Similarity

*Spelling Similarity:
Edit Distance*

Spelling Similarity

- Typos:
 - Brittany Spears → Britney Spears
 - Catherine Hepburn → Katharine Hepburn
 - Reciept → receipt
- Variants in spelling:
 - Theater → theatre

Who Is This?

معمر القذافي

Hints

معمر القذافي




M


Hints



Hints



Hints



Muammar (al-)Gaddafi, or Moamar Khadafi, or ...

Quiz

How many different transliterations can there be?

m
u o
a
m mm
a e
r

el al El Al ø

Q G Gh K Kh
a e u
d dh ddh dhdh th
zz
a
f ff
i y

A Lot!

m u o a m mm a e r	el al El Al ø	Q G Gh K Kh a e u d dh ddh dhdh th zz a f ff i y	
8	x 5	x 360	= 14,400

Edit Operations

- behaviour – behavior (insertion/deletion) (“al”)
- string – spring (substitution) (“k”–”q”)
- sleep – slept (multiple edits)

Levenshtein Method

- Based on dynamic programming
- Insertions, deletions, and substitutions usually all have a cost of 1.

Example

		s	t	r	e	n	g	t	h
	0	1	2	3	4	5	6	7	8
t	1								
r	2								
e	3								
n	4								
d	5								

Recurrence Relation

- Definitions

- $s_1(i)$ – i^{th} character in string s_1
- $s_2(j)$ – j^{th} character in string s_2
- $D(i, j)$ – edit distance between a prefix of s_1 of length i and a prefix of s_2 of length j
- $t(i, j)$ – cost of aligning the i^{th} character in string s_1 with the j^{th} character in string s_2

- Recursive dependencies

$$D(i, 0) = i$$

$$D(0, j) = j$$

$$D(i, j) = \min \begin{bmatrix} D(i-1, j) + 1 \\ D(i, j-1) + 1 \\ D(i-1, j-1) + t(i, j) \end{bmatrix}$$

- Simple edit distance:

$$t(i, j) = 0 \text{ if } s_1(i) = s_2(j)$$

$$t(i, j) = 1, \text{ otherwise}$$

Example

		s	t	r	e	n	g	t	h
	0	1	2	3	4	5	6	7	8
t	1	1							
r	2								
e	3								
n	4								
d	5								

Example

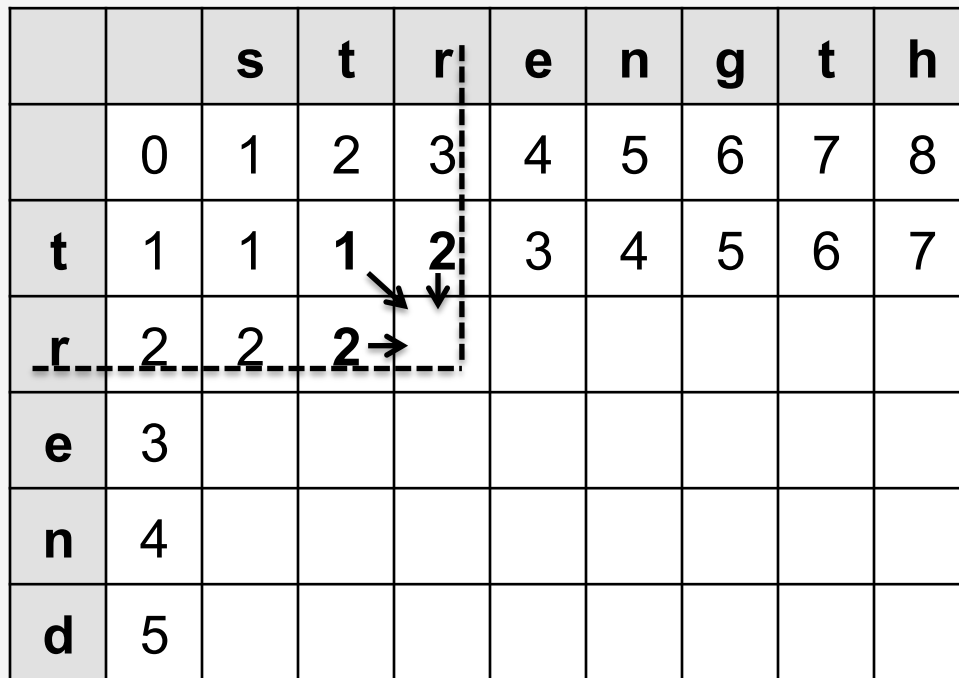
		s	t	r	e	n	g	t	h
	0	1	2	3	4	5	6	7	8
t	1	1	1						
r	2								
e	3								
n	4								
d	5								

Example

		s	t	r	e	n	g	t	h
	0	1	2	3	4	5	6	7	8
t	1	1	1	2	3	4	5	6	7
r	2	2	2						
e	3								
n	4								
d	5								

Example

		s	t	r	e	n	g	t	h
	0	1	2	3	4	5	6	7	8
t	1	1	1	2	3	4	5	6	7
r	2	2	2	2					
e	3								
n	4								
d	5								



Example

		s	t	r	e	n	g	t	h
	0	1	2	3	4	5	6	7	8
t	1	1	1	2	3	4	5	6	7
r	2	2	2	1	2	3	4	5	6
e	3	3	3	2	1	2	3	4	5
n	4	4	4	3	2	1	2	3	4
d	5	5	5	4	3	2	2	3	4

Edit Transcript

		s	t	r	e	n	g	t	h
	0	1	2	3	4	5	6	7	8
t	1	1	1	2	3	4	5	6	7
r	2	2	2	1	2	3	4	5	6
e	3	3	3	2	1	2	3	4	5
n	4	4	4	3	2	1	2	3	4
d	5	5	5	4	3	2	2	3	4

Other Costs

- Damerau modification
 - Swaps of two adjacent characters also have a cost of 1
 - E.g., $\text{Lev}(\text{"cats"}, \text{"cast"}) = 2$,
 $\text{Dam}(\text{"cats"}, \text{"cast"}) = 1$

Quiz

- Some distance functions can be more specialized.
- Why do you think that the edit distances for these pairs are as follows?
 - $\text{Dist}(\text{"sit clown"}, \text{"sit down"}) = 1$
 - $\text{Dist}(\text{"qeather"}, \text{"weather"}) = 1$, *but* $\text{Dist}(\text{"leather"}, \text{"weather"}) = 2$

Quiz Answers

- $\text{Dist}(\text{"sit down"}, \text{"sit clown"})$ is lower in this example because we want to model the type of errors common with optical character recognition (OCR)
- $\text{Dist}(\text{"qeather"}, \text{"weather"}) < \text{Dist}(\text{"leather"}, \text{"weather"})$ because we want to model spelling errors introduced by "fat fingers" (clicking on an adjacent key on the keyboard)



Quiz: Guess the Language

AACCTGCGGAAGGATCATTACCGAGTGCGGGTCCTTTGGGCCCCAACCTCCCATCCGTGTCTATTGTACCC
TGTTGCTTCGGCGGGCCCCGCCGCTTGTCGGCCGCCGGGGGGGCGCCTCTGCCCCCGGGCCCGTGCCCGC
CGGAGACCCCAACACGAACACTGTCTGAAAGCGTGCAGTCTGAGTTGATTGAATGCAATCAGTTAAACT
TTCAACAATGGATCTCTTGGTTCCGGC

Quiz Answer


- This is a genetic sequence (nucleotides AGCT)

>U03518 Aspergillus awamori internal transcribed spacer 1 (ITS1)

AACCTGCGGAAGGATCATTACCGAGTGCGGGTCCTTTGGGCCCCAACCTCCCATCCGTGTCTATTGTACCC
TGTTGCTTCGGCGGGCCCCGCCGCTTGTCGGCCGCCGGGGGGGGCGCCTCTGCCCCCGGGCCCCGTGCCCGC
CGGAGACCCCAACACGAACACTGTCTGAAAGCGTGCAGTCTGAGTTGATTGAATGCAATCAGTTAAACT
TTCAACAATGGATCTCTTGGTTCCGGC

Other Uses of Edit Distance

- In biology, similar methods are used for aligning non-textual sequences
 - Nucleotide sequences, e.g., GTTCGTGATGGAGCG, where A=adenine, C=cytosine, G=guanine, T=thymine, U=uracil, “-”=gap of any length, N=either one of ACGTU, etc.
 - Amino acid sequences, e.g., FMELSEDGIEMAGSTGVI, where A=alanine, C=cystine, D=aspartate, E=glutamate, F=phenylalanine, Q=glutamine, Z=either glutamate or glutamine, X=“any”, etc. The costs of alignment are determined empirically and reflect evolutionary divergence between protein sequences. For example, aligning V (valine) and I (isoleucine) is lower-cost than aligning V and H (histidine).



External URLs

- Levenshtein demo
 - <http://www.let.rug.nl/~kleiweg/lev/>
- Biological sequence alignment
 - http://www.bioinformatics.org/sms2/pairwise_align_dna.html
 - <http://www.sequence-alignment.com/sequence-alignment-software.html>
 - <http://www.ebi.ac.uk/Tools/msa/clustalw2/>
 - <http://www.animalgenome.org/bioinfo/resources/manuals/seqformats>

NACLO Problem

- “Nok–Nok”, NACLO 2009 problem by Eugene Fink:
 - <http://www.naclo.cs.cmu.edu/problems2009/N2009–B.pdf>

Solution to the NACLO Problem

- “Nok–Nok”
 - <http://www.naclo.cs.cmu.edu/problems2009/N2009–BS.pdf>

NACLO Problem

- “The Lost Tram”, NACLO 2007 problem by Boris Iomdin:
 - <http://www.naclo.cs.cmu.edu/problems2007/N2007-F.pdf>

Solution to the NACLO problem

- “The Lost Tram”
 - <http://www.naclo.cs.cmu.edu/problems2007/N2007-FS.pdf>

NLP