

NLP

Introduction to NLP

Hidden Markov Models (cont'd)

Observation Likelihood

- Given multiple HMMs
 - e.g., for different languages
 - Which one is the most likely to have generated the observation sequence
- Naïve solution
 - try all possible state sequences

Forward Algorithm

- Dynamic programming method
 - Computing a forward trellis that encodes all possible state paths.
 - Based on the Markov assumption that the probability of being in any state at a given time point only depends on the probabilities of being in all states at the previous time point

HMM Learning

- Supervised
 - Training sequences are labeled
- Unsupervised
 - Training sequences are unlabeled
 - Known number of states
- Semi-supervised
 - Some training sequences are labeled

Supervised HMM Learning

- Estimate the static transition probabilities using MLE

$$a_{ij} = \frac{\text{Count}(q_t = s_i, q_{t+1} = s_j)}{\text{Count}(q_t = s_i)}$$

- Estimate the observation probabilities using MLE

$$b_j(k) = \frac{\text{Count}(q_i = s_j, o_i = v_k)}{\text{Count}(q_i = s_j)}$$

- Use smoothing

Unsupervised HMM Training

- Given:
 - observation sequences
- Goal:
 - build the HMM
- Use EM (Expectation Maximization) methods
 - forward-backward (Baum-Welch) algorithm
 - Baum-Welch finds an approximate solution for $P(O|\mu)$

Outline of Baum-Welch

- Algorithm
 - Randomly set the parameters of the HMM
 - Until the parameters converge repeat:
 - E step – determine the probability of the various state sequences for generating the observations
 - M step – reestimate the parameters based on these probabilities
- Notes
 - the algorithm guarantees that at each iteration the likelihood of the data $P(O|\mu)$ increases
 - it can be stopped at any point and give a partial solution
 - it converges to a local maximum

NLP